



Summary

- Architecture of Database Systems
- Transaction Management
- Modern Database Technology
- Data Warehouses and OLAP
- Data Mining
- Big Data Analytics


Storage Layout



Query Execution Plans



Data Processing Models




U+H
Universität Hamburg
DER FORSCHUNG | DER LEHRE | DER BILDUNG

Effect of Engine Optimizations


Simple query on small data (<300kB)

SELECT * FROM testsequence WHERE Sequence_Count<3;



Complex query on slightly more data (~20MB)

SELECT count(*) FROM (SELECT 1 FROM testsequence, authors WHERE Sequence_Count < 3 AND testsequence.FOB_ID = authors.LOCCODE GROUP BY authors.author) foo;



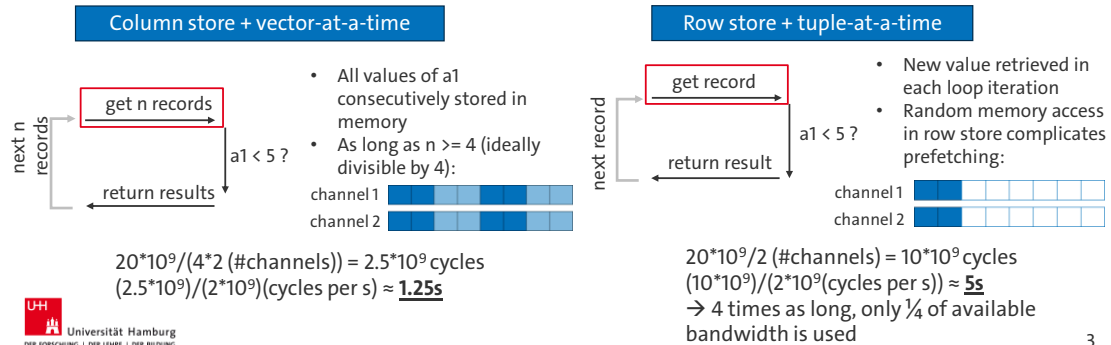
SQLite vs DuckDB

Recap Data Access Time

SELECT a1 FROM t WHERE a1 < 5;

a1 ∈ SMALLINT
20*10⁹ records
2 memory channels @ 2GHz & 64 bit

Minimum time required for data transfer from memory to CPU for the following scenarios:



Thinking about what happens on the hardware side is always a useful tool for optimizing performance

More interesting aspects of this example:

- The example assumes that the data is stored in a compact format in the column store, i.e. only the space needed for the data format is used, not the size of a whole processor word. If this is not the case, the required time increases → This is the absolute „minimum“ if the memory is used optimally
- The time required for the comparison itself can most likely be masked because it is done faster than the retrieval of the new record(s) which happens at the same time
- Writing data is slower than reading it → There is a break-even point when operator-fusion (if there was more than 1 operator) becomes more beneficial than vector-at-a-time processing, but it is different depending on the data type (i.e. its size), the (physical) operators, and the hardware


Optimizations by the User: Indexes

- Primary or Secondary Index (additional access path)
- Simple index or multilevel index
- Storage structure:
 - Tree-structured, e.g. B-Tree
 - Sequential, e.g. Sequential lists
 - Scattered, e.g. Hash regions
- Access method:
 - Key comparison
 - Key transformation, e.g. hashing

Sequential List

blau	
braun	
grün	TID(123, 3) TID(125, 4)
lila	
rot	
schwarz	

Binary search tree



Recap of indexes in 1st DIS-lecture

Optimizations by the User: (Materialized) Views

Store (materialize) the view Create a view called CheapFood


```
CREATE MATERIALIZED VIEW CheapFood AS SELECT Meal FROM MensaMeals WHERE Price < 5;
```

Refresh the view

```
REFRESH MATERIALIZED VIEW CheapFood;
```



- Refresh the view after updates in your base data
- Materialization not supported by all database systems



5

Optimizations in RDBMS

Reusability of queries and query results

- Queries and Subqueries (Views) can be stored and referenced → Nicer queries
- The result of views can be stored → Higher performance for frequently used queries and remote data

Types of Materialized Views

- Materialized join view
- Materialized aggregate view
- Materialized aggregate-join view

Derivability of queries

- Query Q may be replaced by Q'
 - Q' uses the contents of mat. view V
 - Q' delivers a semantically equivalent results to Q
 - Q' is called compensation query
- Identification of derivability is NP-hard

- Predicate P_M und P_Q : $P_Q \subseteq P_M$

- General undecidability of predicate logic
- Reduction to elemental predicates after standardization in disjunctive normal form
- Projected attributes are not considered (but would be necessary for decidability)
- More about computability theory in general: “Theoretische Informatik - kurzgefasst” by Uwe Schöning, or in any textbook about the theory of informatics
- More about predicate logic: “Logik und Logikprogrammierung” by Steffen Hölldobler (chapter 4.8.4 is about the undecidability of predicate logic)

Exercise Optimizations by the User

Query 1

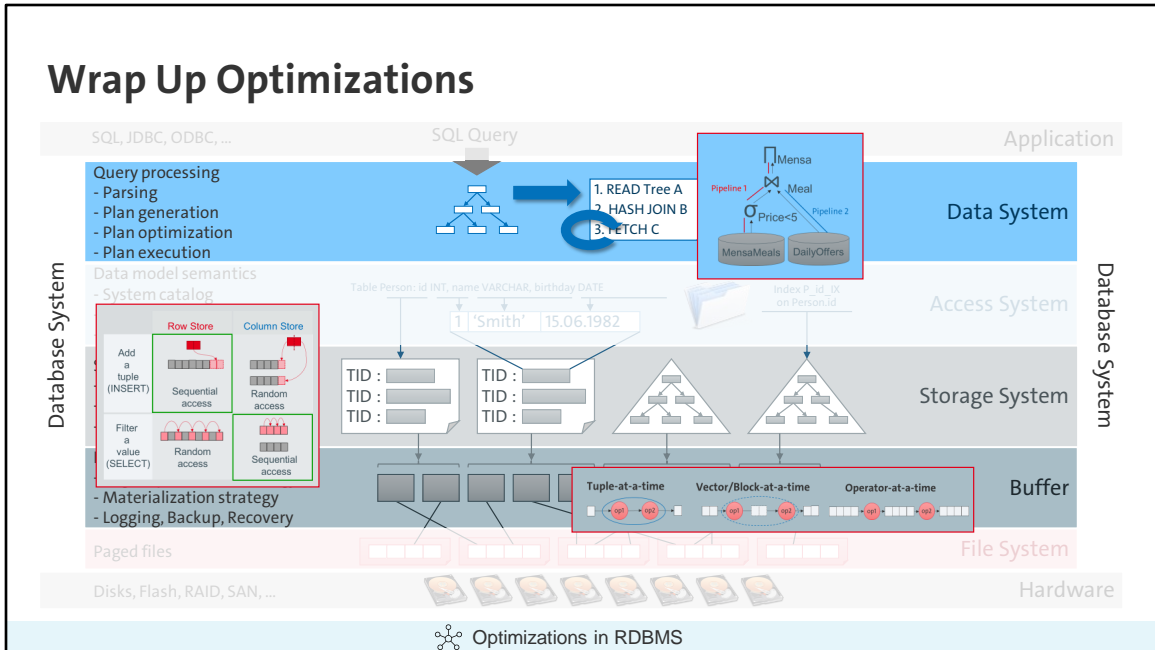
```
SELECT count(*) FROM testsequence, authors WHERE testsequence.Sequence_Count < 3 AND  
testsequence.PDB_ID = authors.IDCODE;
```

Query 2

```
SELECT authors.institute,authors.name FROM institute, testsequence, authors WHERE  
testsequence.Sequence_Count < 3 AND testsequence.PDB_ID = authors.IDCODE AND  
institute.name = authors.institute;
```

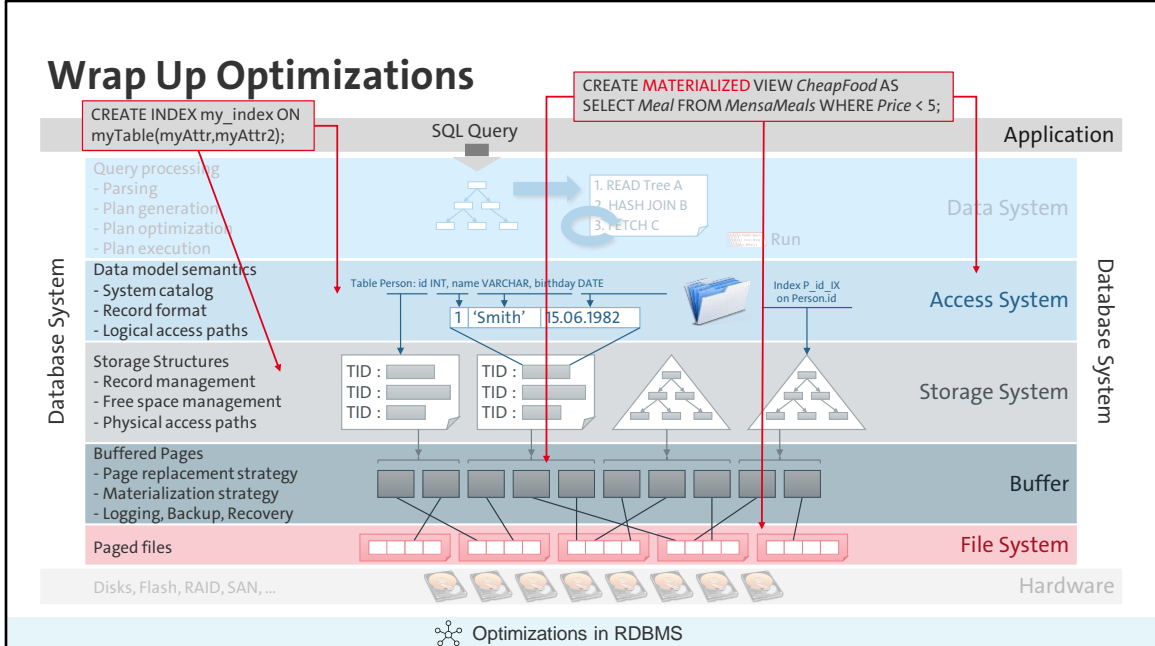
Which optimizations can a user apply to increase the query performance?

Wrap Up Optimizations



- There are countless more optimizations that can be done, e.g. different physical operators, intelligent prefetching,...
- The Data System takes into account the features of the Access System for its optimization, e.g. available indexes
- Column Stores usually provide a better performance for analytical queries than row stores

Wrap Up Optimizations



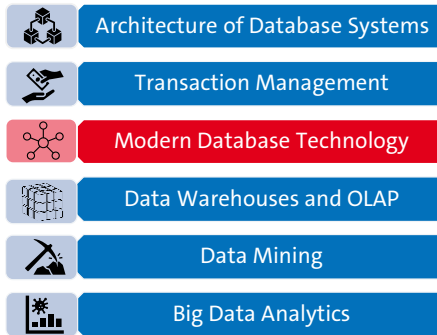
- The optimizations a user can apply (which are not changing/enabling/disabling a feature of the DBS as a whole) are usually triggered by SQL queries
- They change the behavior in the lower layers
 - A materialized view can change the files on disc, materializes a query result that might otherwise stay in the buffer, and adds an additional access path to data resulting from the query in question
 - Indexes add additional access paths (stored in the storage system and used by the access system)
- Some other optimizations, e.g. query rewriting, might be useful but their effect depends on the used DBS

Survey



<https://evasys-online.uni-hamburg.de/evasys/online.php?pswd=J3Q6H>


Course Outline



- Distributed Systems
- Optimizations in RDBMS
- **NoSQL**

Big Data Management

2000's – Today: Big Data



"data comes first, schema comes second" (or never)



Need to revisit Data Management

The number of Vs depends on who you ask, it started with just Variety, Volume, and Velocity

Challenges for Data Management

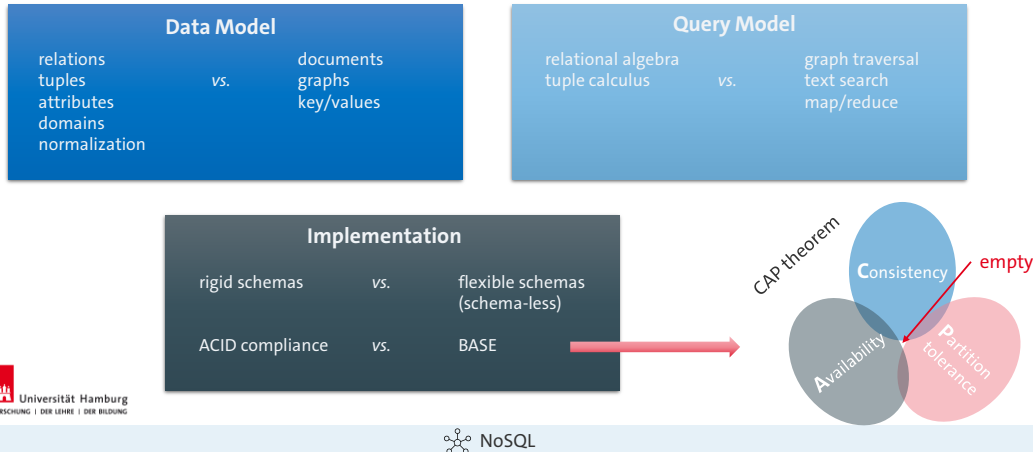
- Volume: From GB's to PB's ...
- Variety: Tables, Text, Images, ...
- Velocity: Sensors continuously generate data
- Value: More complex analytics (from SQL to deep learning)
- Veracity: Data quality, data bias, ...

Big data scenarios try to avoid a priori data modelling

- Use case unknown, data sources unknown, source schemas unknown/permanently changing, etc.
- NoSQL as movement to reflect an agile "working with existing data", instead of "reflecting a real world by a schema/logical model"
- If a logical data model is created prior to a NoSQL implementation, the Logical Model and the Physical Model may differ – even substantially, especially due to the extreme de-normalization and flattening that occurs within NoSQL

Non-Relational Databases

Most NoSQL data systems diverge from standard relational systems



ACID → Atomicity, Consistency, Isolation, Durability

BASE: Availability over consistency

- **B**asically available: The database is available for changes, even if these changes are applied later, i.e. it is not necessarily in a consistent state when the user accesses it
- **S**oft state: Data can have a temporary state, e.g. when multiple applications change it at the same time
- **E**ventually consistent: Consistency is reached eventually, but only when (finally) all changes have been applied that were made at the same time

→ BASE scales easier than ACID

CAP theorem:

- In a distributed system, only 2 of the 3 properties Availability, Consistency, and partition tolerance can be reached at the same time
- ACID prioritizes Consistency (CA, CP)
- BASE prioritizes Availability (AP)

Types of Data Systems

Relational Databases (SQL)


- Popular Systems
IBM DB2 PostgreSQL, MS SQL Server, SQLite, MySQL
- NewSQL Database
(Scale-out using distributed processing)
Google Spanner, VoltDB (H-Store), SAP HANA, MemSQL

Non-relational/Not only SQL Databases (NoSQL)

- Key-Value Pair (KVP) Databases
Redis, Riak, RocksDB
- Document Databases
MongoDB, CouchDB
- Wide Column Databases
Hbase, Cassandra, Druid, Vertica
- Graph Databases
Neo4J, Allegro, InfiniteGraph, OrientDB, Virtuoso

NoSQL Databases: Key-Value Stores

Key-Value Stores




Get ... a value by a given key
Put ... a new key-value pair into the database
Delete ... a key and the associated value

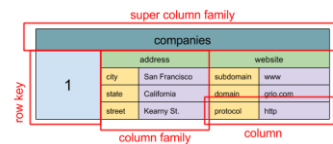
Value is a (semi) structured document

Value is a row or a table

Document Stores



Wide Column Stores



Key-value Stores

- Minimal set of functions:
 - void put(key,value)
 - value = get (key)
 - void delete(key)
- Might implement additional features, e.g. collections, merge operator, purge operator,...
- Does not differentiate between datatypes, i.e. the content of arbitrary documents can be stored (in a binary format)

Wide Column Stores


- Not the same as column stores (see last lecture)
- Image source: <https://vishnuch.tech/4-types-of-nosql-databases>, accessed: 21.02.2022
- Suitable for distributed systems
- Queries can look similar to queries in key-value store, e.g. (Hbase):
 - put 'students','student1','account:name','Alice'
 - put 'students','student1','address:country','GB'
- Queries can also look similar to SQL, e.g. (Hbase):

```
SELECT CONVERT_FROM(row_key, 'UTF8') AS studentid,  
       CONVERT_FROM(students.account.name, 'UTF8') AS name,  
       CONVERT_FROM(students.address.country, 'UTF8') AS country  
FROM students;
```

Document Stores

- Image source: <https://lennilobel.wordpress.com/2015/06/01/relational-databases-vs-nosql-document-databases/>
- Usually supports json and/or xml
- Allows querying individual fields of the supported documents
- Different languages might be supported, e.g. (MongoDB shell):
 - `db.inventory.find({ status: "A" })` (*inventory is a collection*)

NoSQL Databases: Graph DBs



16

NoSQL

Most popular GraphDB: Neo4j

Basis

- Assuming a domain of objects \mathcal{O}
- A graph is a structure (V, E)
- Vertices are objects: $V \subseteq \mathcal{O}$
- Edges are 2-element subset of vertices: $E \subseteq \{\{x, y\} | x, y \in V\}$

Direction

- Edges are ordered pairs of vertices $E \subseteq V \times V$

Multigraph

- Multiple edges per node
- A graph is a structure (V, E, ρ)
- Edges are objects: $E \subseteq \mathcal{O}$
- $\rho: E \rightarrow \{\{x, y\} | x, y \in V\}$ (undirected) / $\rho: E \rightarrow V \times V$ (directed)
is a function assigning to each edge
a 2-element subset of / ordered pair of vertices

Hypergraphs

- Edges can connect more than two vertices
- $G=(V,E)$ with $E\subseteq 2^V$

Labels

- Descriptive type information tagged to the objects
 - Indicate which kind of real world entity an object represents
 - Assuming a domain of labels \mathcal{L}
 - A graph is a structure (V,E,λ)
 - $\lambda:V\rightarrow\mathcal{L}$ (vertex label)
 - $\lambda:E\rightarrow\mathcal{L}$ (edge label)
 - $\lambda:V\cup E\rightarrow\mathcal{L}$ (vertex and edge label)
 - $\lambda:V\rightarrow 2^\mathcal{L}$ (vertex labels)
 - $\lambda:E\rightarrow 2^\mathcal{L}$ (edge labels)
 - $\lambda:V\cup E\rightarrow 2^\mathcal{L}$ (vertex and edge labels)
- is a partial / total function assigning to each
vertex / edge / vertex and edge
a label / a set of labels

Food for thought: (No)SQL

- Which kind of database would be suited for the following use cases in your opinion?
 - A social network
 - The cover text of books and their ISBN number
 - The mapping between students/employees and their transponder number
 - The structure of different molecules (= atoms and their bindings)
 - A lab book
 - The inventory of a library
 - Banking

Yes, multiple systems might be suitable if there is a good explanation