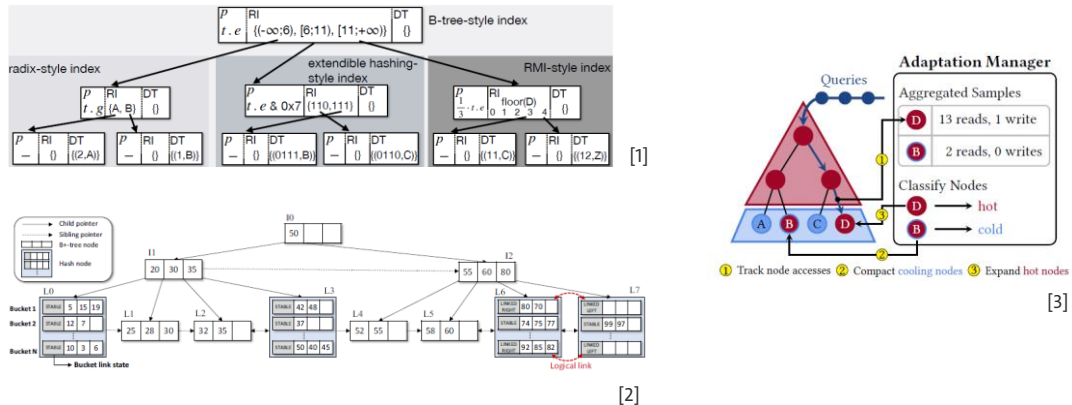


Addendum on Optimizations: Hybrid Indexes



[1] Dittrich, Jens, Joris Nix, and Christian Schön. "The next 50 years in database indexing or: the case for automatically generated index structures." *Proceedings of the VLDB Endowment* 15.3 (2021): 527-540.


[2] Cha, Hokeun, et al. "Blink-hash: An adaptive hybrid index for in-memory time-series databases." *Proceedings of the VLDB Endowment* 16.6 (2023): 1235-1248.

[3] Anneser, Christoph, et al. "Adaptive hybrid indexes." *Proceedings of the 2022 International Conference on Management of Data*. 2022.

Summary

- Architecture of Database Systems
- Transaction Management
- Modern Database Technology
- Data Warehouses and OLAP
- Data Mining
- Big Data Analytics


User optimizations: Indexes and Materialized Views






CREATE **MATERIALIZED VIEW** CheapFood
AS
SELECT Meal FROM MensaMeals
WHERE Price < 5;

REFRESH MATERIALIZED VIEW CheapFood;


Wrap up optimizations



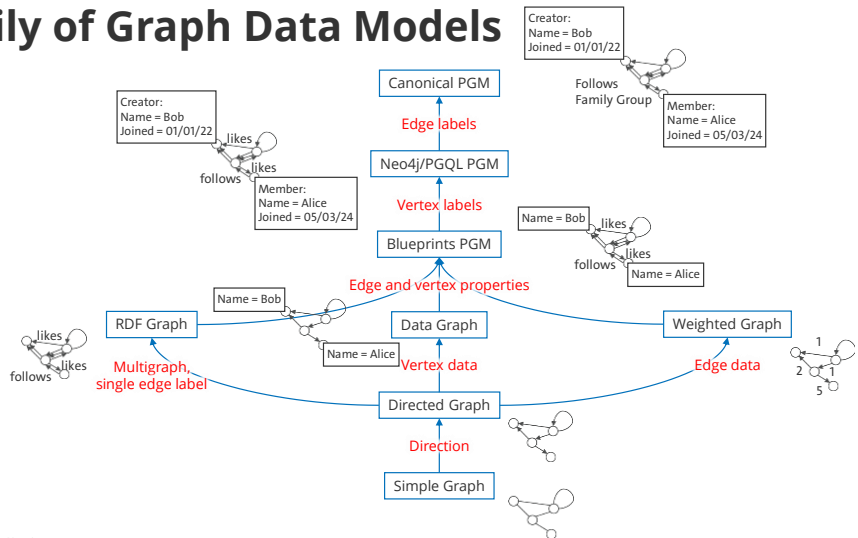
NoSQL Systems



NoSQL DBS: Challenges & Characteristics



Family of Graph Data Models



PGM - Property Graph Model

RDF - Resource Description Framework

→ Stored in triple format (node, edge label, node)

→ Each property must be modeled as a triple Label: Descriptive Type information

PGQL – Property Graph Query Language (tries to look like SQL)

Property/Data - Actual data

Label – Descriptive (type) information

→ Labels are sometimes called “type”, especially in the context of edges

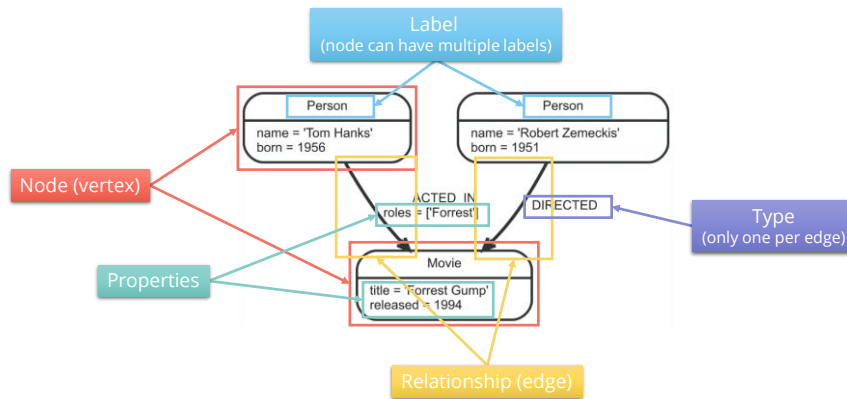
Properties/Data:

- Assuming a domain of data / values \mathcal{D} and a domain of property keys / attributes \mathcal{K}
- A graph is a structure (V, E, ν)
- $\nu: V \rightarrow \mathcal{D}$ (vertex data)
- $\nu: E \rightarrow \mathcal{D}$ (edge data)
- $\nu: V \cup E \rightarrow \mathcal{D}$ (vertex and edge data)
- $\nu: V \times \mathcal{K} \rightarrow \mathcal{D}$ (vertex properties)
- $\nu: E \times \mathcal{K} \rightarrow \mathcal{D}$ (edge properties)

$\nu: (V \cup E) \times \mathcal{K} \rightarrow \mathcal{D}$ (vertex and edge properties)
is a partial function assigning data / values to each
vertex / edge / vertex and edge

Neo4j PGM


[<http://neo4j.com/docs/developer-manual/current/introduction/#graphdb-concepts>]



Property Graph Modelling

Entity with Properties


- Entity relationship model



- Schema typically implicit, i.e. given with instances

```
(ja:Person { name: 'Jane Austen', yearOfBirth: 1775 })
```

- Instance



Person <id>: 175 yearOfBirth: 1775 name: Jane Austen

Vertex id is system generated in Neo4j


Cypher (Neo4j) uses the keyword CREATE to create a new vertex, e.g.

```
CREATE (ja:Person { name: 'Jane Austen', yearOfBirth: 1775 })
```


Property Graph Modelling

Relationships (N:M)

- Entity relationship model



- Vertices
 - `(ja:Person { name: 'Jane Austen', yearOfBirth: 1775 })`
 - `(wb:Person { name: 'William Blake', yearOfBirth: 1757 })`
 - `(lo:City {name: 'London'})`
 - `(ba:City {name: 'Bath'})`
 - `(fe:City {name: 'Felpham'})`
- Edges
 - `(ja)-[:livedIn {since: 1775}]->(lo)`
 - `(ja)-[:livedIn {since: 1800}]->(ba)`
 - `(wb)-[:livedIn {since: 1757}]->(lo)`
 - `(wb)-[:livedIn {since: 1800}]->(fe)`




An edge can be created with the CREATE keyword in Cypher, just like vertices, e.g.
CREATE (ja)-[:livedIn {since: 1775}]->(lo)


Property Graph Modelling

n-ary relationships (with $n > 2$)

- Entity relationship model



```
(ja:Person { name: 'Jane Austen', yearOfBirth: 1775 })  
(wh:Publisher { name: 'Whitehall' })  
(sas:Book {title: 'Sense and Sensibility' })  
(pub:Publication {year: 1811 })  
(pub)-[:author]->(ja)  
(pub)-[:book]->(wh)  
(pub)-[:publisher]->(sas)
```



Not directly expressible since edge connect exactly two vertices, not more, not less

→ Solution: Reification of relationship

Property Graph Modelling

Multivalued properties



```
(wb:Person {name: 'William Blake', hadOccupation: ['Poet','Painter','Printmaker']})
```



Person <id>: 183 **hadOccupation:** Poet,Painter,Printmaker **name:** William Blake

Expressibility depends on datatype system of specific system in use

→ Possible in Neo4j (see slide)

→ If not expressible: Reification of values of multivalued property

Multiple labels can be connected with a : or a & in Neo4j, e.g. (wb:Person&Artist {name: 'William Blake', hadOccupation: ['Poet','Painter','Printmaker']})


Exercise Neo4j

- Your fridge is stocked with cheese, butter, milk, and pizza
- There are also 4 tomatoes and two bars of chocolate in the kitchen
- Each item has an expiration date and a quantity: cheese (15 May, 10), butter (01 June, 1), milk (11 May, 2), Pizza (25 Aug, 5)
- You have a recipe for luxury frozen pizza that is called “PiDeLuxe”. It requires a slice of cheese and a tomato per frozen pizza. You want to prepare one Pizza “PiDeLuxe” today.


How can this be modeled using the ER model?
How could it be expressed with Neo4j?

Representations: Adjacency Matrix

Directed graph with edge labels




Directed graph with edge labels



Formalisms to define graph operations



Reachability

$$\begin{pmatrix} 0 & 1 & 1 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{pmatrix} \cdot \begin{pmatrix} 0 & 1 & 1 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 1 \end{pmatrix}$$



Graph Isomorphism

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} 0 & 1 & 1 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{pmatrix} \cdot \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}^{-1} = \begin{pmatrix} 0 & 0 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

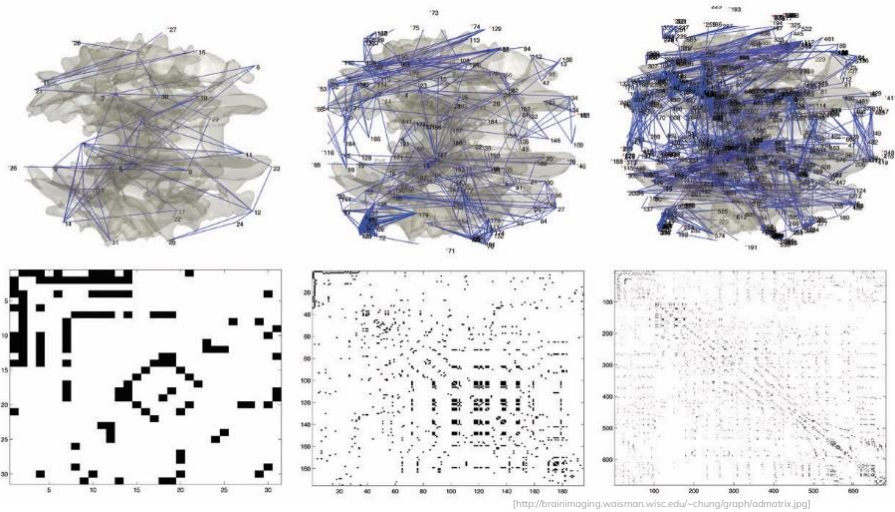


Reachability: $A^2 = A \cdot A$ shows all paths of length 2 between nodes of graph represented by A


Graph isomorphism: Graphs represented by A_1 and A_2 are isomorphic if there exists a permutation matrix P such that $PA_1P^{-1}=A_2$

Note: Most definitions work only for unlabeled graphs!

Adjacency Matrix



14



Adjacency matrices are often sparse
→ Compression

Compress Sparse Row (CSR)


$$\begin{array}{c} 0 \\ 1 \\ 2 \\ 3 \end{array} \begin{pmatrix} a & & & \\ & b & & \\ d & i & f & \\ & h & g & \end{pmatrix}$$

Position for row # in other two arrays:

#	0	1	2	3
Row position array:	0	2	4	7


#	0	1	2	3	4	5	6	7	8
Column index array:	0	2	1	3	0	1	2	1	2
Cell value array:	a	c	b	e	d	i	f	h	g

15



Also works column-wise → Compress Sparse Column (CSC)

Exercise CSR



What does the adjacency matrix for this graph look like?
Apply Compress Sparse Row.

Representations: Adjacency List

Source vertex with outgoing edges...

	0	1	2	3
0	a		c	
1		b		e
2	d	i	f	
3		h	g	

...without edge labels

0 -> (0,2)
1 -> (1,3)
2 -> (0,1,2)
3 -> (1,2)

...with edge labels

0 -> ([0,a],[2,c])
1 -> ([1,b],[3,e])
2 -> ([0,d],[1,i],[2,f])
3 -> ([1,h],[2,g])

...with edge properties

0 -> ([0,a,(weight=4)],[2,c,(weight=3)])
1 -> ([1,b,(weight=3)],[3,e,(weight=2)])
2 -> ([0,d,(weight=5)],[1,i,(weight=2)],...)
3 -> ([1,h,(weight=9)],[2,g,(weight=7)])

The same!

Almost the same!

Compressed
Sparse Row

Compressed
Sparse Column

source-oriented

0	2	4	7
a	c	b	e
d	i	f	h
g			

target-oriented

0	2	5	8
a	d	b	i
h	c	f	g
e			

Representations: Relational Representations I


Triple Table: A table with three columns: subject, predicate, object

Subject	Predicate	Object
<http://www.w3.org/People/EM/contact#me>	<http://www.w3.org/2000/10/swap/pim/contact#fullName>	"Eric Miller"
<http://www.w3.org/People/EM/contact#me>	<http://www.w3.org/2000/10/swap/pim/contact#mailbox>	<mailto:e.miller123(at)example>
<http://www.w3.org/People/EM/contact#me>	<http://www.w3.org/2000/10/swap/pim/contact#personalTitle>	"Dr."
<http://www.w3.org/People/EM/contact#me>	<http://www.w3.org/1999/02/22-rdf-syntax-ns#type>	<http://www.w3.org/2000/10/swap/pim/contact#Person>

Vertex and Edge Table: Two universal tables, one for vertices, one for edges

ID	Type	Color	Name	RAM	Nationality	Source	Target	Type	Rating
1	Product	black	"Apple iPad MC707LL/A"	64 GB		1	7	in	
...
4	Category		"Cell Phones & Accessories"			5	4	part of	
5	Category		"Phones"			7	6	part of	
...

19



Triple Table

- Generic, simple, straightforward approach
- Can be built on relational storage engines
- Add indexing for faster lookups
- Queries need to operate on all triples (no data portioning based on schema information)
- High storage redundancy

Vertex and Edge Table

- Requires efficient handling of NULL values
 - Given for instance in column stores
- Queries need to operate on all vertices/edges (no data portioning based on schema information)

Further reading

Daniel J. Abadi et al. Scalable Semantic Web Data Management Using Vertical Partitioning. VLDB 2007

Representations: Relational Representations II

Clustered Property Table: Groups properties that tend to be defined together in a table

Property Table

Subj.	Type	Title	copyright
ID1	BookType	"XYZ"	"2001"
ID2	CDType	"ABC"	"1985"
ID3	BookType	"MNP"	NULL
ID4	DVDType	"DEF"	NULL
ID5	CDType	"GHI"	"1995"
ID6	BookType	NULL	"2004"

Left-Over Triples

Subj.	Prop.	Obj.
ID1	author	"Fox, Joe"
ID2	artist	"Orr, Tim"
ID2	language	"French"
ID3	language	"English"

Property-Class Table: Cluster similar sets of subjects together in the same table

Class: BookType

Subj.	Title	Author	copyright
ID1	"XYZ"	"Fox, Joe"	"2001"
ID3	"MNP"	NULL	NULL
ID6	NULL	NULL	"2004"


Class: CDType

Subj.	Title	Artist	copyright
ID2	"ABC"	"Orr, Tim"	"1985"
ID5	"GHI"	NULL	"1995"

Left-Over Triples

Subj.	Prop.	Obj.
ID2	language	"French"
ID3	language	"English"
ID4	type	DVDType
ID4	title	"DEF"

20



Clustered Property Table

- Example:
 - Properties: type, title, and copyright date
 - Table stores all triples with one of these properties
- Multiple property tables with different clusters of properties may be created
- One particular property may only appear in at most one property table
- Triples that not fit any property table are stored in a left-over triple table
- Multivalued attributes are problematic
- Queries that have unspecified property are problematic
- Introduce NULL values

Property-Class Table

- Exploits the type property of subjects
- A property may exist in multiple property-class tables
- A subject may also exist in multiple tables
- Multivalued attributes are problematic
- Queries that do not select on class type are problematic
- Introduce NULL values

Property Table Approaches

Triple Table

Subj.	Prop.	Obj.
ID1	type	BookType
ID1	title	"XYZ"
ID1	author	"Fox, Joe"
ID1	copyright	"2001"
ID2	type	CDType
ID2	title	"ABC"
ID2	artist	"Orr, Tim"
ID2	copyright	"1985"
ID2	language	"French"
ID3	type	BookType
ID3	title	"MNO"
ID3	language	"English"
ID4	type	DVDType
ID4	title	"DEF"
ID5	type	CDType
ID5	title	"GHI"
ID5	copyright	"1995"
ID6	type	BookType
ID6	copyright	"2004"

(Clustered) property table

Property Table

Subj.	Type	Title	copyright
ID1	BookType	"XYZ"	"2001"
ID2	CDType	"ABC"	"1985"
ID3	BookType	"MNP"	NULL
ID4	DVDType	"DEF"	NULL
ID5	CDType	"GHI"	"1995"
ID6	BookType	NULL	"2004"

Left-Over Triples

Subj.	Prop.	Obj.
ID1	author	"Fox, Joe"
ID2	artist	"Orr, Tim"
ID2	language	"French"
ID3	language	"English"

Property-Class Table

Class: BookType

Subj.	Title	Author	copyright
ID1	"XYZ"	"Fox, Joe"	"2001"
ID3	"MNP"	NULL	NULL
ID6	NULL	NULL	"2004"

Class: CDType

Subj.	Title	Artist	copyright
ID2	"ABC"	"Orr, Tim"	"1985"
ID5	"GHI"	NULL	"1995"

Left-Over Triples

Subj.	Prop.	Obj.
ID2	language	"French"
ID3	language	"English"
ID4	type	DVDType
ID4	title	"DEF"


Reduce numbers of subject-subject self joins necessary to reconstruct entities

Property Graph Queries in SQL

```
CREATE PROPERTY GRAPH bank_sql_pg
  VERTEX TABLES (
    bank_accounts
    KEY (id)
    LABEL account
    PROPERTIES (owner, balance)
  )
  EDGE TABLES (
    bank_txns
    KEY (txn_id)
    SOURCE KEY (from_acct_id) REFERENCES bank_accounts (id)
    DESTINATION KEY (to_acct_id) REFERENCES bank_accounts (id)
    LABEL transfer
    PROPERTIES (amount, reference)
  );
```


```
SELECT * FROM GRAPH_TABLE (bank_sql_pg MATCH
(a IS account WHERE a.id = 42) -[e IS transfer]->(b IS account)
COLUMNS (a.id AS acc_a, e.amount AS amount, b.id AS acc_b)
);
```

Example



transfer
txn_id = 12347 (unique)
amount = 5.00
reference = 'kitty'

transfer
txn_id = 12345 (unique)
amount = 15.00
reference = 'pizza order'



This is supposed to be an arrow

ACC_A	AMOUNT	ACC_B
42	15.00	15
42	5.00	15

NoSQL

- PGQs were recently adopted in the SQL standard (2023): ISO/IEC DIS 9075-16
- First implementation in Oracle version 23.2:
<https://docs.oracle.com/en/database/oracle/property-graph/23.2/spgdg/sql-property-graphs.html>
- Only because PGQs are supported by SQL, there is no guarantee that the DBS implementing it is optimized for graph processing!**